

# Distributed Data Processing Environments

José Orlando Pereira

Departamento de Informática  
Universidade do Minho

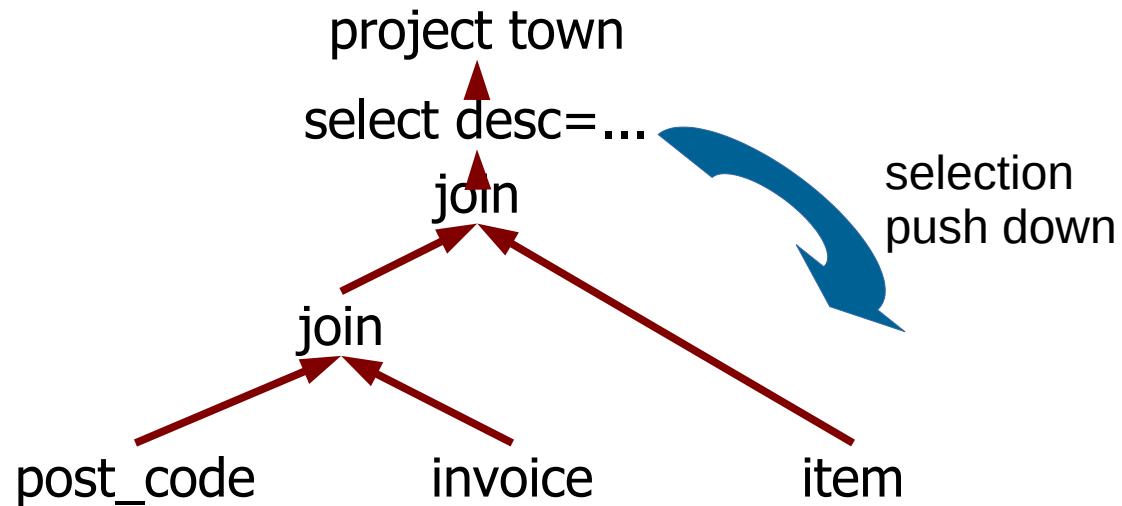


# Roadmap

- What physical operators exist for each logical operation?
- How are physical operators implemented and composed?
- How is the best plan found?

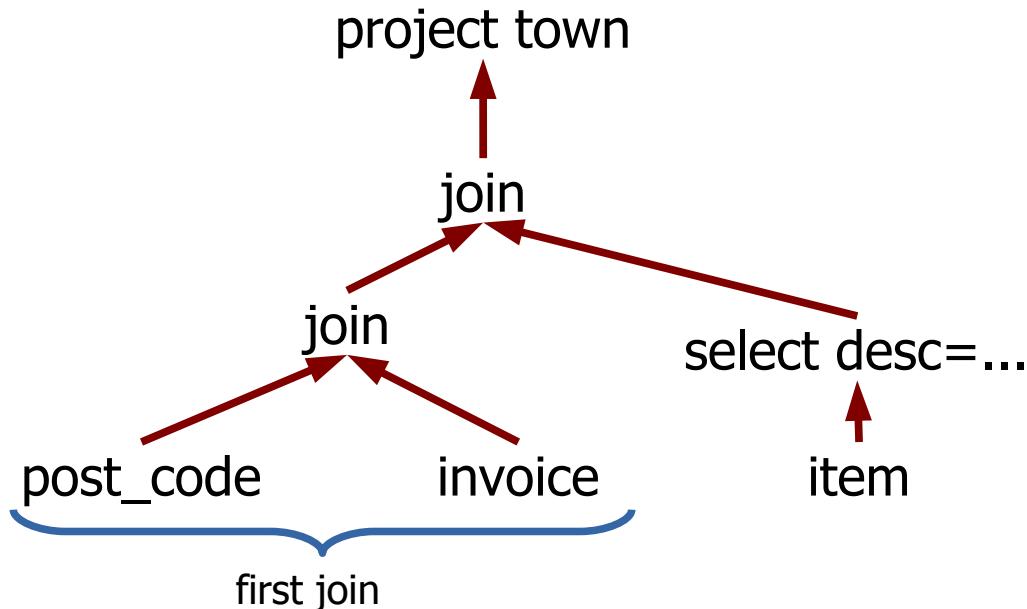
# Optimization

```
select town from  
post_code natural join invoice  
natural join item where desc=...
```



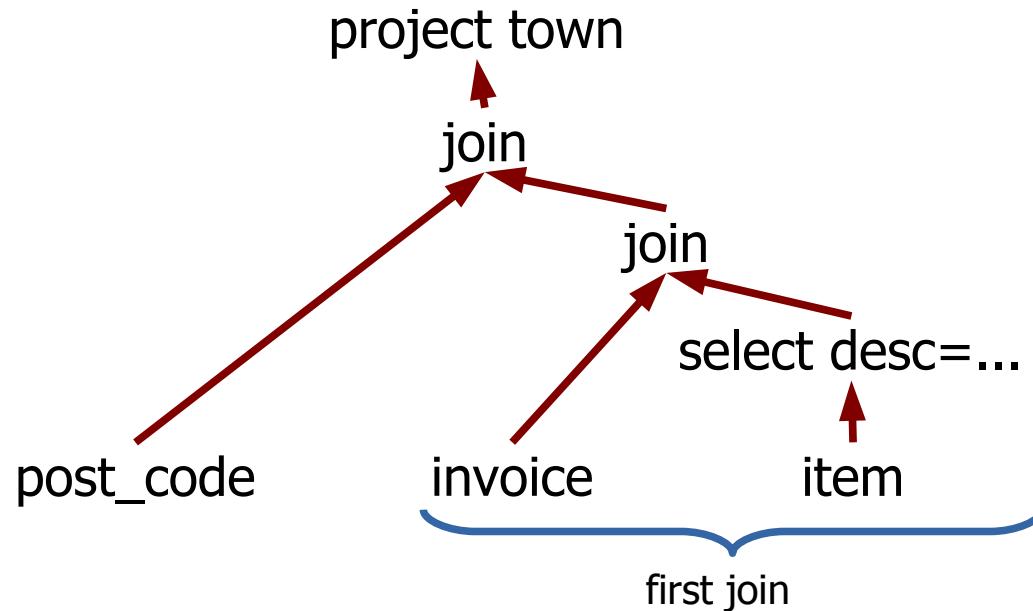
# Optimization

```
select town from  
post_code natural join invoice  
natural join item where desc=...
```



# Optimization

```
select town from  
post_code natural join invoice  
natural join item where desc=...
```



# Roadmap

- How to estimate the cost of a plan?
- How to find alternative plans?

# Cost estimation

- Tradeoff between:
  - Actually executing the query and measuring what resources it consumes and how long it takes to execute
    - vs
  - **Estimate that can be computed quickly**
    - vs
  - Considering that all queries have the same cost
- We don't want to know what is the actual cost
- We want to know which alternative costs less
  - Use a model that monotonously approximates real cost

# Example: Simple sequential scan

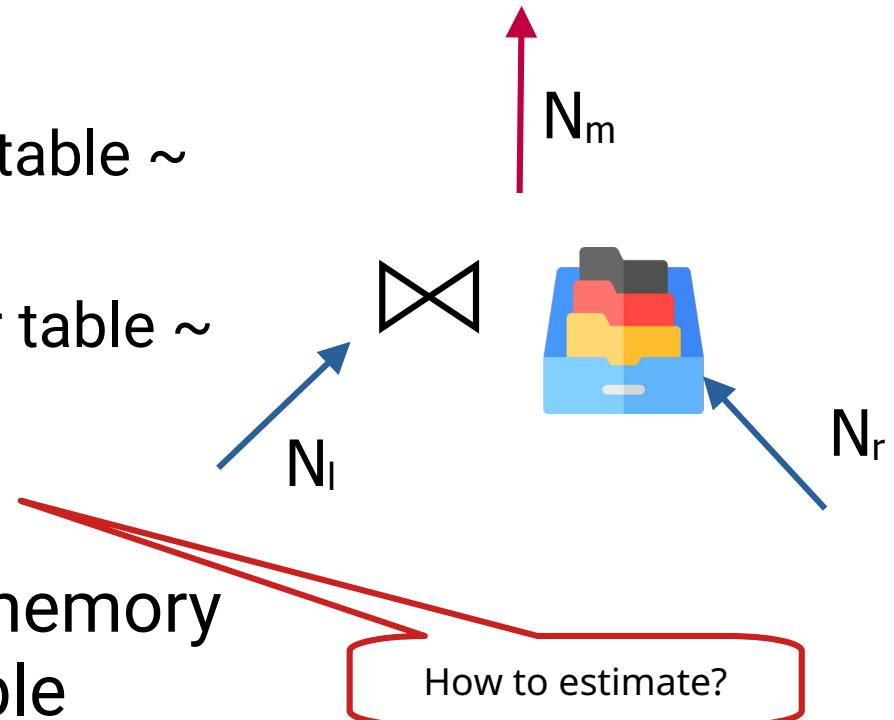
- Assumption of main cost factor:
  - Number of disk block I/O operations
- Cost model as a simple equation:
  - $C = C_0 + C_1 N_{\text{blocks}}$

Known from the size  
of the file

How to find them?

# Example: Hash Join

- CPU cost:
  - Building hash table from inner table ~  $N_r$  input rows
  - Checking each row in the outer table ~  $N_i$  input rows
  - Creating the resulting  $N_m$  rows
- Very high cost if not enough memory to hold  $N_i$  rows in the hash table



# Coefficients

- Cost coefficients depend on the actual hardware, software, and even workloads
- Can be estimated by profiling simple workloads
- Can be tuned by the DBA

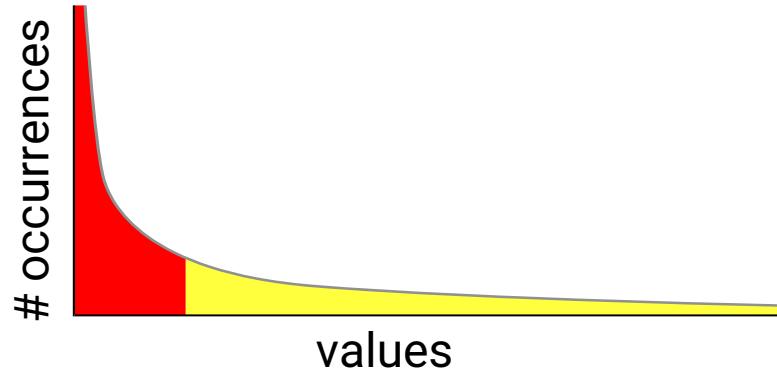
# Cardinality - Selection

- Assumption of uniform distribution
- Know #distinct
  - Expected copies of each tuple:
    - #rows / #distinct

# Cardinality - Selection

- Real data are usually not uniformly distributed:

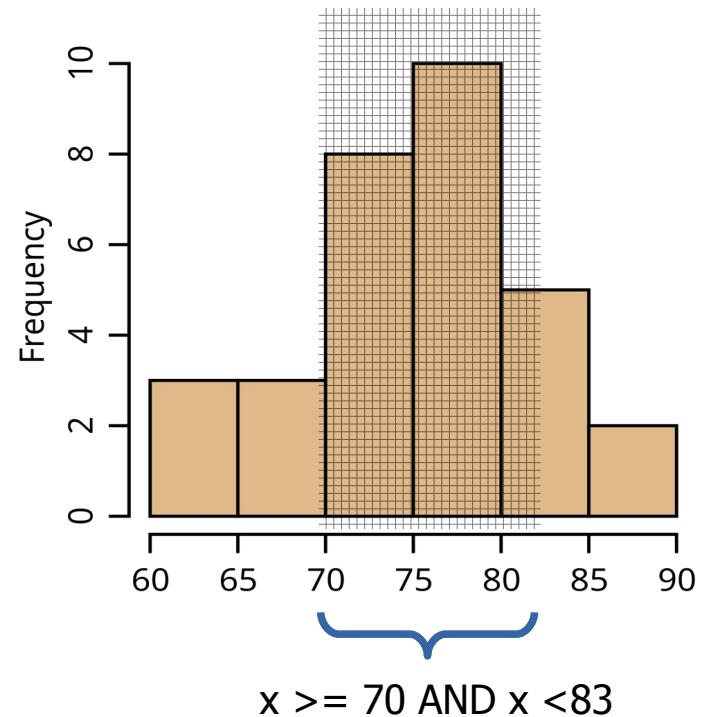
- 80/20 rule
  - Power law



- Know most popular and #occurrences
- Compute estimated #occurrences of others as uniformly distributed

# Cardinality - Selection

- Data often have complex multimodal distributions
- Know histogram:
  - % of occurrences in each interval or
  - interval for fixed % of occurrences
- Compute #occurrences as uniformly distributed within each interval



# Cardinality - Conjunction

- Filter conditions are often the conjunction of conditions on different columns
- Statistics on multiple columns are expensive:
  - Multidimensional
  - Many possible combinations
- Assume independently distributed values in different columns:
  - selectivity = #selected / #total rows
  - $\text{selectivity}(a \wedge b) = \text{selectivity}(a) * \text{selectivity}(b)$
- Idem for disjunction

# Cardinality - Join

- Cardinality for cross-product of A and B:
  - $\# \text{rows from A} \times \# \text{rows from B}$
- Cardinality for join, first attempt:
  - Align buckets of histograms for A and B
  - Estimate the cardinality of each matching bucket as a cross-product:
    - Not good, as there are non-matching values
      - e.g. A has even numbers, B has odd numbers  $\rightarrow$  no match!

# Cardinality - Join

- Assume that one relation A has all values (containment)
- Cardinality for join, second attempt:
  - Align buckets of histograms for A and B
  - Estimate the cardinality of each matching bucket:
    - Each of #rows in B matches ( $\# \text{rows in A} / \# \text{distinct in A}$ )
      - Because A has all the values
    - Estimate as  $\# \text{rows in B} \times \# \text{rows in A} / \# \text{distinct in A}$
  - Generalize for the case where B has all values:
    - $\# \text{rows in A} \times \# \text{rows in B} / \max(\# \text{distinct in A}, \# \text{distinct in B})$

# Summary

- Tradeoff between complexity (data and computation) and accuracy
- Many other techniques:
  - More statistics
  - Heuristics
  - Hinting
  - Sampling of query
  - Machine learning
  - ...

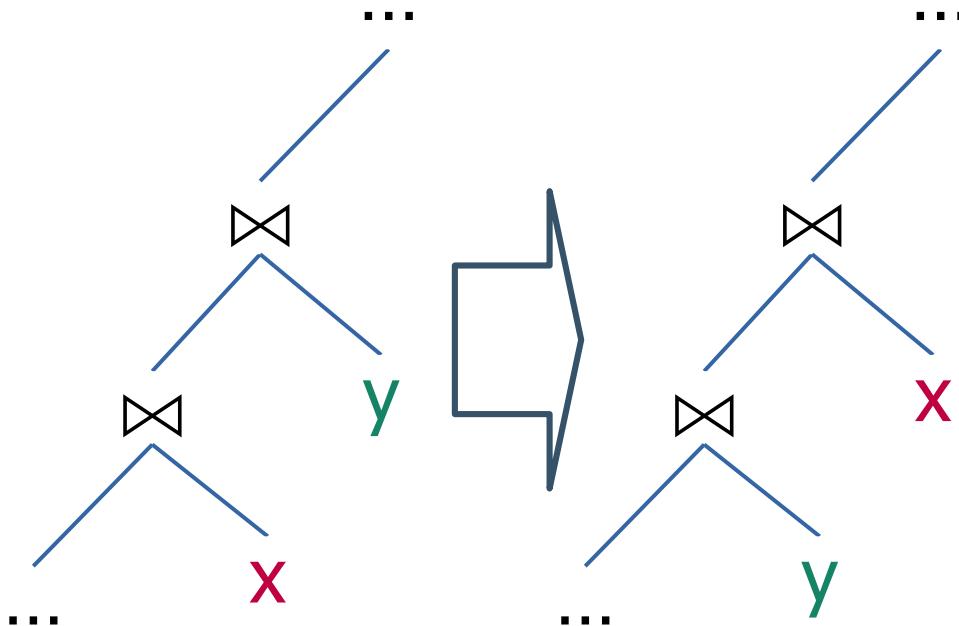
# Roadmap

- How to estimate the cost of a plan?
- How to find alternative plans?

# Search space

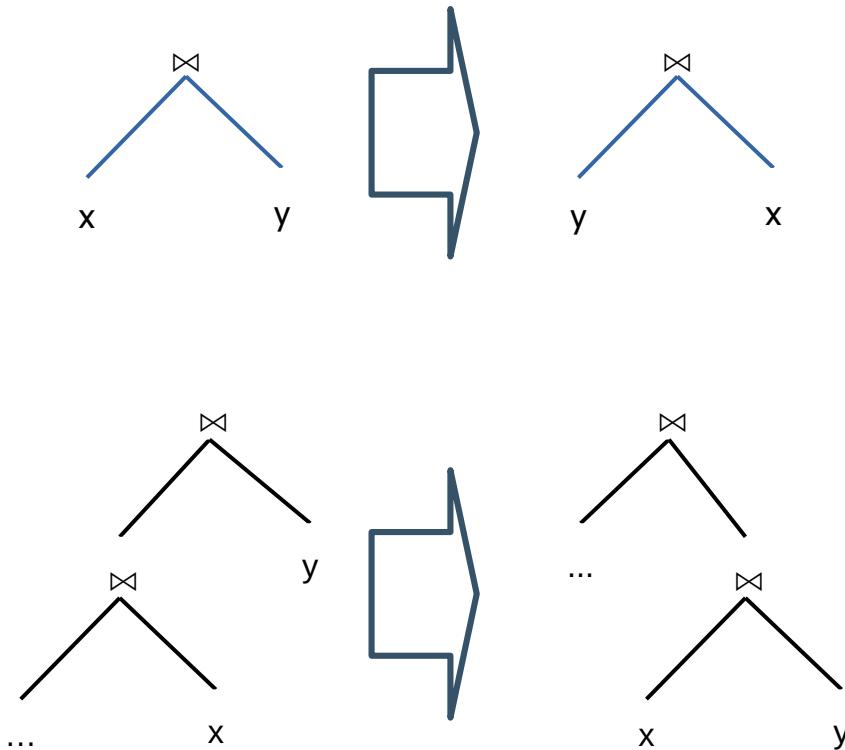
- The set of possible alternative plans (search space) is determined by a set of rules
  - Equivalent relational algebra expressions
  - Physical implementation of single operators or plan fragments
  - Enforcing physical properties
- The set of rules is the main configuration point for extensible query optimizers

# A simple rule for Join order



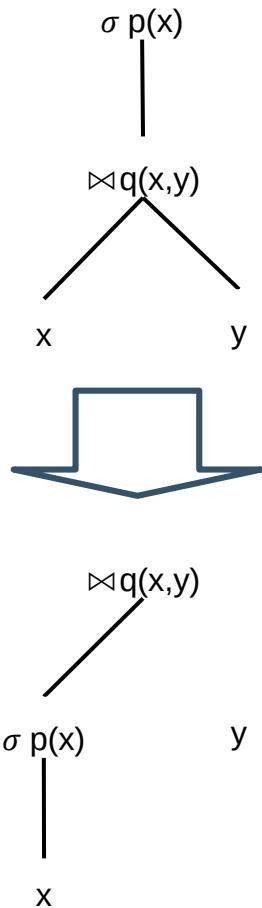
- Inner join is commutative and associative
$$(\dots \bowtie x) \bowtie y =$$
$$\dots \bowtie (x \bowtie y) =$$
$$\dots \bowtie (y \bowtie x) =$$
$$(\dots \bowtie y) \bowtie x$$
- This allows exploring all left-deep trees
  - $n!$  permutations
- Does not consider bushy trees

# Join commutativity and associativity



- Separately considering commutativity and associativity rules
- Allows exploring all possible join expressions:
  - Including bushy trees

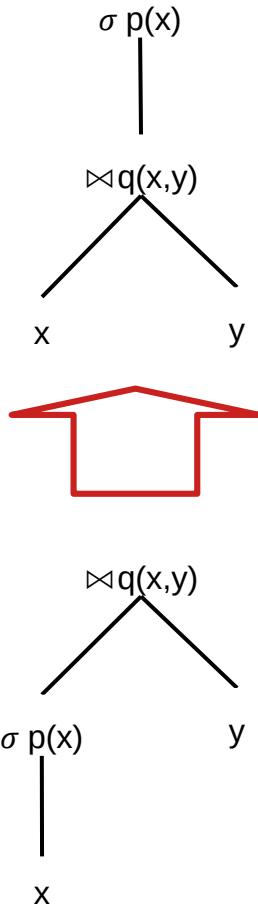
# Selection push-down



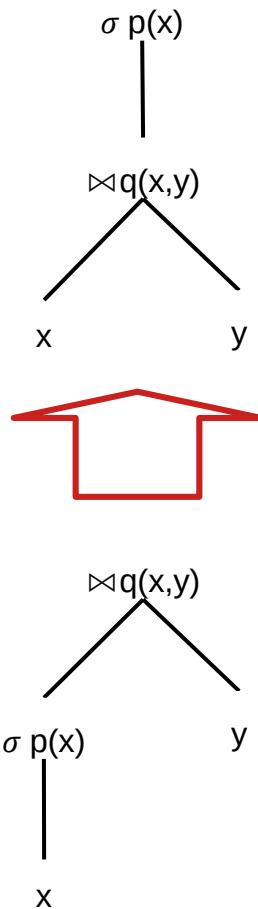
- Possible when the predicate involves only one of the branches
  - Otherwise, merge it in the join condition!
- If the predicate is highly selective, reduces the amount of work in join
- Selectivity is the key criteria for join ordering!!!

# Selection pull-up

- Always possible, but...
- Does it ever make sense?!?!

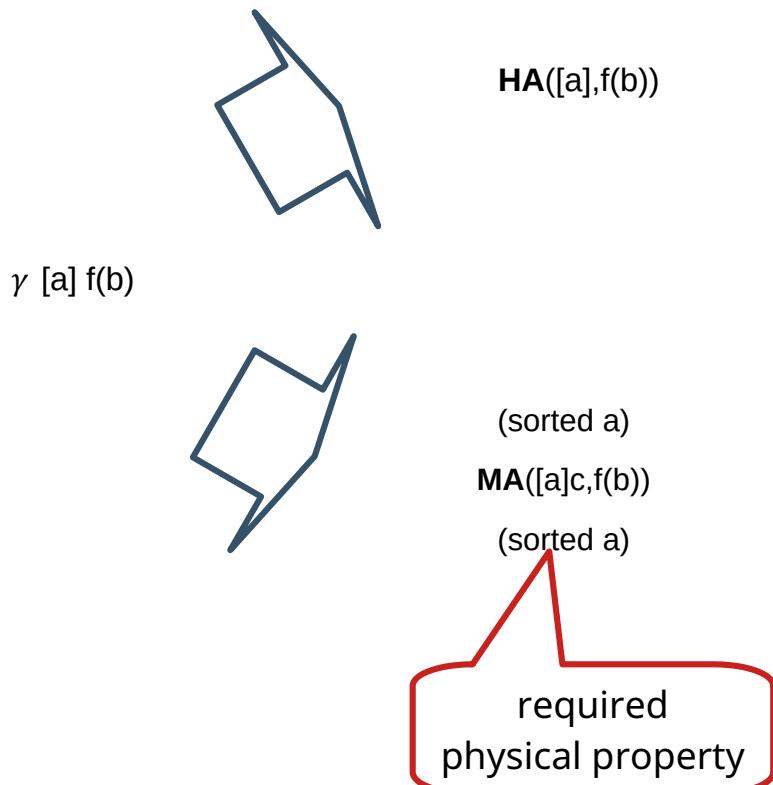


# Selection pull-up



- Assume that  $q$  is highly selective
- Assume that  $p$  is **very costly** to execute:
  - e.g. call into an LLM to check if “sentiment” on a textual column is “positive”
- Can be useful!

# Grouping and aggregation implementation

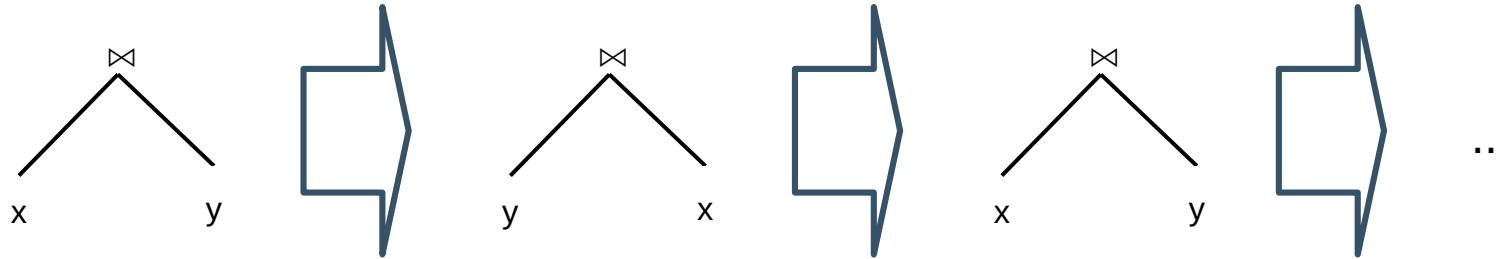


- Two-pass algorithms for grouping (and join) depend on sorted input
  - Expressed as a required “physical property”
- Sorting is preserved

# Search space

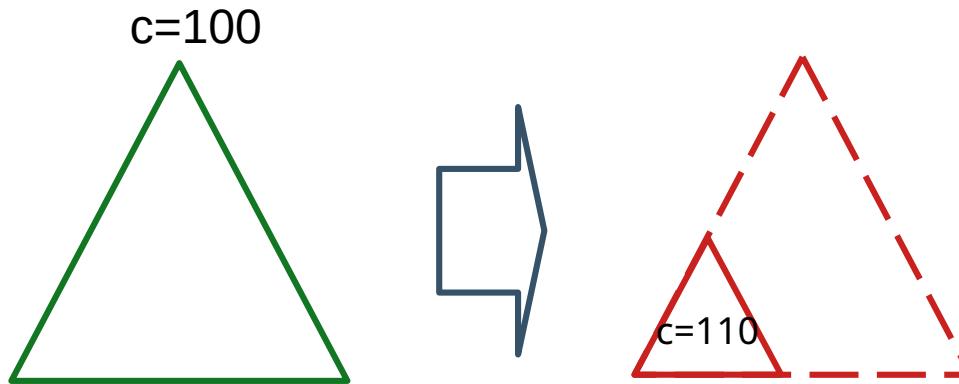
- Rule sets that produce a larger search space:
  - More likely to contain the optimum → faster execution
  - More work to evaluate all alternatives → faster planning
- Query execution is the sum of planning and execution
- Best overall performance:
  - Depends on workload
  - Found as a compromise of planning and execution

# Search algorithm



- Repeated application of rules can result in an infinite loop
- Solution:
  - Remember all plans to check for repetitions

# Search algorithm



- The cost of a sub-plan of an alternative being explored may be greater than the total cost of the best known alternative
- Solution:
  - Prune search based on current best estimate

# Search algorithm

- Some sequences of rule applications tend to converge faster to the optimum plan
  - The sooner that we get a “good” estimate, the more alternatives that can be pruned
- Solution:
  - Order available rules by their heuristic promise

# Conclusions

- The set of possible alternative execution plans for common analytical operations is extremely large
- There is no greedy algorithm that can find the optimum
- The optimum depends on current data, defeating manual optimization efforts
- The best option is to use a declarative data processing system such as a SQL query engine