

# Distributed Data Processing Environments

José Orlando Pereira

Departamento de Informática  
Universidade do Minho




# Roadmap

- Problem too big to be efficiently stored and processed by a single server
  - Distributed-parallel data processing
- Problem too complex to be expressed as a single step and/or with a single tool
  - DAG orquestration

# Map Reduce


Input data set



- Data is split in chunks that can be handled separately
- Done by the framework, but might need help from the programmer

# Map Reduce


Input data set



- Input data items are mapped to zero or more key/value pairs
- Map function is provided by the programmer
  - Can deal with arbitrary and unstructured data formats (e.g., plain text)

# Map Reduce


Input data set



- Key-value pairs are grouped by key
- Done by the framework

# Map Reduce

Input data set




Output data set

- Sequences of values for each key are reduced to 0 or more output data items
- Reduce function is provided by the programmer

# Map Reduce

Input data set



- When reduce is associative, values for the same key from the same mapper can be combined
- Lessens the data to be grouped

# Selection and projection

- Example:  
select x, y+1 from ... where z = ...;
- Can be performed by the Map stage:
  - Return computed key/value for those items that match


# Grouped aggregates

- Example:  
    select x,sum(y) from ... group by x;
- In the Map stage:
  - Return Key=x and Value=y
- Optionally use a Combiner stage
- In the Reduce stage:
  - Iterate over values and return Key,sum(Values)

# Map/Broadcast join

- Example:  
    select a.y,b.z from A join B on a.x=b.x;
- Assumptions:
  - One data set (B) is small
  - No assumption on number of occurrences
- Before Map, cache B in all workers
- In the Map stage:
  - Lookup a.x in B, getting b.z
  - Return a.y and b.z

# Map/Broadcast join



# Shuffle join

- Example:

select a.y,b.z from A join B on a.x=b.x;

- Assumptions:

- Both input data sets are large
  - Small number of occurrences of each key


- In the Map stage:

- For A: Return Key=x and Value=(LEFT,y)
  - For B: Return Key=x and Value=(RIGHT,z)

- In the Reduce stage:

- Collect all (LEFT,...) and (RIGHT,...) pairs
  - Return all combinations

# Shuffle




# Sort

- Example:  
select x,y from ... order by y;
- In the Map stage:
  - Return Key=y and Value=x
- Use one identity reducer task

# Chaining


- Multiple MapReduce jobs can be chained to compose operations
- This is necessary when sorting and grouping need to be done on multiple criteria
  - e.g., join on x, group by y

# MR: Efficiency limitations



Need to avoid writing to disk so much!

# MR: Usability limitations



Need to freely compose stages!

# Distributed dataflow



- Generalized MapReduce:
  - Map, shuffle, reduce, and persistence stages
  - Can be arbitrarily composed
- Efficiency improved with:
  - Caching
  - Query optimization and code generation

# Example


```
public class GroupByTest {
    public static void main(String[] args) throws Exception {
        int numMappers = 100;
        int numKVPairs = 10000;
        int valSize = 1000;
        int numReducers = 36;

        SparkConf conf = new SparkConf().setMaster("local").setAppName("GroupBy Test");
        JavaSparkContext sc = new JavaSparkContext(conf);

        List<Integer> data = IntStream.range(0, numMappers).boxed().collect(Collectors.toList());
        JavaPairRDD<Integer, byte[]> pairs1 = sc.parallelize(data, numMappers)
            .flatMapToPair(p -> {
                Random ranGen = new Random();
                Stream<Tuple2<Integer, byte[]>> arr1 = IntStream.range(0, numKVPairs).mapToObj(i -> {
                    byte[] byteArr = new byte[valSize];
                    ranGen.nextBytes(byteArr);
                    return new Tuple2<>(ranGen.nextInt(), byteArr);
                });
                return arr1.iterator();
            });
        long count1 = pairs1.count();
        long count2 = pairs1.groupByKey(numReducers).count();

        System.out.println(count1+" "+count2);
    }
}
```

# Example



# Caching example


```
public class GroupByTest {
    public static void main(String[] args) throws Exception {
        int numMappers = 100;
        int numKVPairs = 10000;
        int valSize = 1000;
        int numReducers = 36;

        SparkConf conf = new SparkConf().setMaster("local").setAppName("GroupBy Test");
        JavaSparkContext sc = new JavaSparkContext(conf);

        List<Integer> data = IntStream.range(0, numMappers).boxed().collect(Collectors.toList());
        JavaPairRDD<Integer, byte[]> pairs1 = sc.parallelize(data, numMappers)
            .flatMapToPair(p -> {
                Random ranGen = new Random();
                Stream<Tuple2<Integer, byte[]>> arr1 = IntStream.range(0, numKVPairs).mapToObj(i -> {
                    byte[] byteArr = new byte[valSize];
                    ranGen.nextBytes(byteArr);
                    return new Tuple2<>(ranGen.nextInt(), byteArr);
                });
                return arr1.iterator();
            })
            .cache(); // Store interim result for both jobs
        long count1 = pairs1.count();
        long count2 = pairs1.groupByKey(numReducers).count();

        System.out.println(count1+" "+count2);
    }
}
```


# Query processing



# Example

- Consider the following code:

```
Dataset<Row> ds = spark.sql("select town from " +  
    "post_code natural join invoice " +  
    "natural join item where desc='stuff'");  
  
ds.show();  
  
ds.explain(true);
```




Shows how the query is  
processed in the Spark SQL pipeline

# Example

- Parsed logical plan:


```
'Project ['town]
+- 'Filter ('desc = stuff)
  +- 'Join NaturalJoin(Inner)
    :- 'Join NaturalJoin(Inner)
      :  :- 'UnresolvedRelation [post_code], [], false
      :  +- 'UnresolvedRelation [invoice], [], false
    +- 'UnresolvedRelation [item], [], false
```



# Example

- Analyzed logical plan:

```
town: string
Project [town#59]
+- Filter (desc#39 = stuff)
  +- Project [item#17, code#58, town#59, invoice#16, desc#39]
    +- Join Inner, (item#17 = item#38)
      :- Project [code#58, town#59, invoice#16, item#17]
      :  +- Join Inner, (code#58 = code#18)
      :  :- SubqueryAlias post_code
      :  :  +- Relation[code#58,town#59] csv
      :  +- SubqueryAlias invoice
      :     +- Relation[invoice#16,item#17,code#18] csv
    +- SubqueryAlias item
      +- Relation[item#38,desc#39] csv
```




# Example

- Optimized logical plan:

```
Project [town#59]
+- Join Inner, (item#17 = item#38)
  :- Project [town#59, item#17]
  :  +- Join Inner, (code#58 = code#18)
  :    :- Filter isnotnull(code#58)
  :    :  +- Relation[code#58,town#59] csv
  :    +- Project [item#17, code#18]
  :      +- Filter (isnotnull(code#18) AND isnotnull(item#17))
  :        +- Relation[invoice#16,item#17,code#18] csv
  +- Project [item#38]
    +- Filter ((isnotnull(desc#39) AND (desc#39 = stuff)) AND isnotnull(
      +- Relation[item#38,desc#39] csv
```

- Using `ds.explain("cost")` shows additional statistics




# Code generation

- Converts a physical plan into actual code that can be executed
- Simple code generation:
  - Produces sequence of generic RDD transformations
- Whole stage code generation:
  - Produces custom RDD transformation combining all operations in each stage
  - Resulting code is very different from manually written code...

# Example

- Physical plan:

```
*(3) Project [town#59]
+- *(3) BroadcastHashJoin [item#17], [item#38], Inner, BuildRig
  :- *(3) Project [town#59, item#17]
  :  +- *(3) BroadcastHashJoin [code#58], [code#18], Inner, Bu
  :    :- BroadcastExchange HashedRelationBroadcastMode(List(
  :      :    +- *(1) Filter isnotnull(code#58)
  :      :      +- FileScan csv [code#58,town#59] Batched: false
  :      :      +- FileScan csv [item#17,code#18] Batched: false, D
  +- BroadcastExchange HashedRelationBroadcastMode(List(input[
  +- *(2) Project [item#38]
    +- *(2) Filter ((isnotnull(desc#39) AND (desc#39 = stu
      +- FileScan csv [item#38,desc#39] Batched: false, D
```



- Using `ds.explain("codegen")` shows final Java code

# Summary

- Distributed-parallel processing is applicable for extremely large datasets
- SQL compilation is extended to use distributed operators:
  - Data exchange as a new dimension for optimization